Machine Learning – Clustering

Professor Widom's Instructional Odyssey

www.professorwidom.org

Stanford University

Google

Google Cloud Platform

Association for Computing Machinery

Very Large Data Bases Endowment Inc.

Instabase

amazon web services™

Big Data Tools and Techniques

- Basic Data Manipulation and Analysis

  Performing well-defined computations or asking well-defined questions ("queries")

- Data Mining

  Looking for patterns in data

- Machine Learning

  Using data to build models and make predictions

- Data Visualization

  Graphical depiction of data

- Data Collection and Preparation

Machine Learning

Using data to build models and make predictions

Supervised machine learning

- Set of labeled examples to learn from: training data
- Develop model from training data
- Use model to make predictions about new data

Unsupervised machine learning

- Unlabeled data, look for patterns or structure (similar to data mining)

Clustering

Like classification, data items consist of values for a set of features (numeric or categorical)

- Medical patients

  Feature values: age, gender, symptom1-severity, symptom2-severity, test-result1, test-result2

- Web pages

  Feature values: URL domain, length, #images, heading$_1$, heading$_2$, …, heading$_n$

- Products

  Feature values: category, name, size, weight, price

Clustering

Like classification, data items consist of values for a set of features (numeric or categorical)

Unlike classification, there is no label

- Medical patients

  Feature values: age, g
  symptom2-severity, test-result1, test-result2

- Web pages

  Feature values: URL domain, length, #images, heading_1, heading_2, …, heading_n

- Products

  Feature values: category, name, size, weight, price

Clustering

Like K-nearest neighbors, for any pair of data items i_1 and i_2, from their feature values can compute distance function: *distance*(i_1,i_2)

Example:

Features - gender, profession, age, income, postal-code

person$_1$ = (male, teacher, 47, \$25K, 94305)
person$_2$ = (female, teacher, 43, \$28K, 94309)

distance(person$_1$, person$_2$)

distance() can be defined as inverse of *similarity*()

Clustering

GOAL: Given a set of data items, partition them into groups (= clusters) so that items within groups are close to each other based on distance function

- ➢ Sometimes number of clusters is pre-specified
- ➢ Typically clusters need not be same size

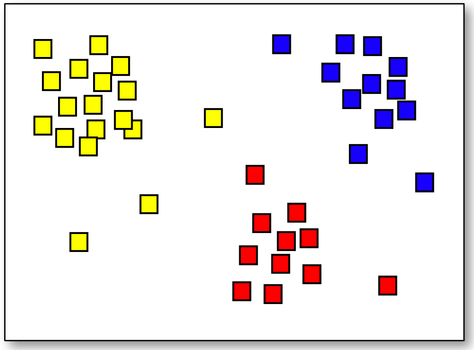

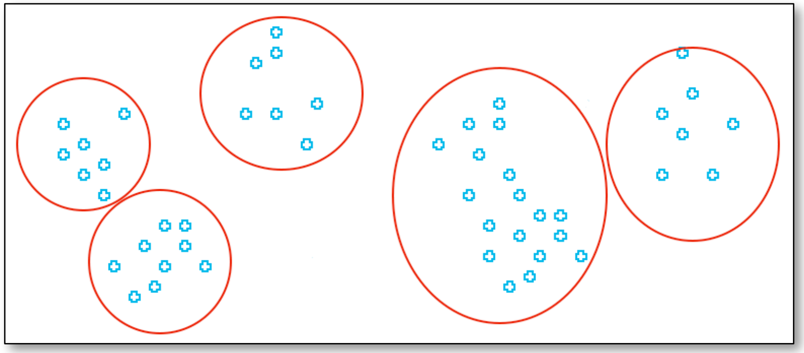

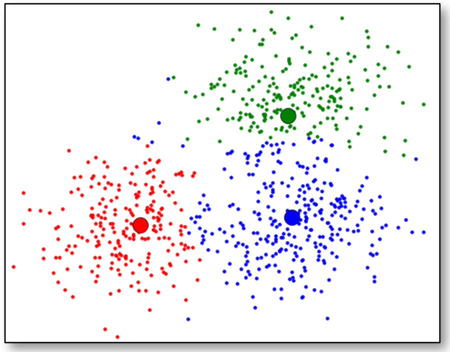

Some Uses for Clustering

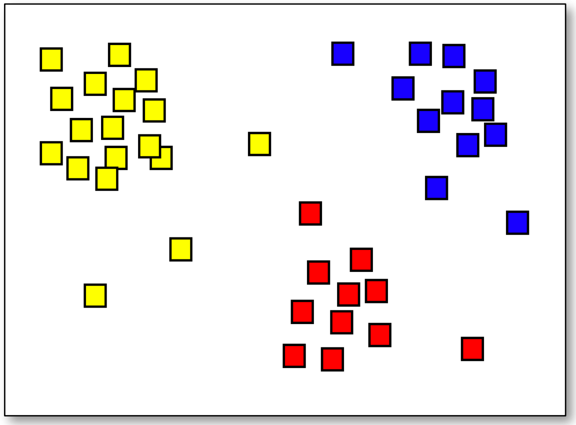

- Classification!
  - Assign labels to clusters
  - New data items get the label of their cluster
- Identify similar items
  - For substitutes or recommendations
  - For de-duplication
- Anomaly (outlier) detection
  - Items that are far from any cluster

K-Means Clustering

Reminder: for any pair of data items i_1 and i_2 have *distance*(i_1,i_2)

For a group of items, the mean value (centroid) of the group is the item *i* (in the group or not) that minimizes the sum of *distance*(*i*,*i’*) for all *i’* in the group

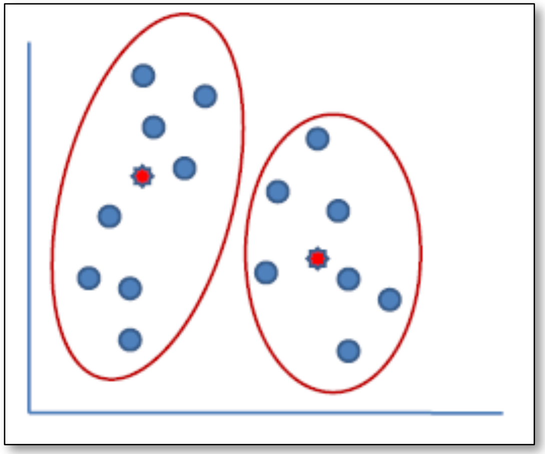

K-Means Clustering

For a group of items, the mean value (centroid) of the group is the item *i* (in the group or not) that minimizes the sum of *distance*(*i*,*i'*) for all *i'* in the group

- Error for each item: distance *d* from the mean for its group; squared error is d^2
- Error for the entire clustering: sum of squared errors (SSE)

Remind you of anything?

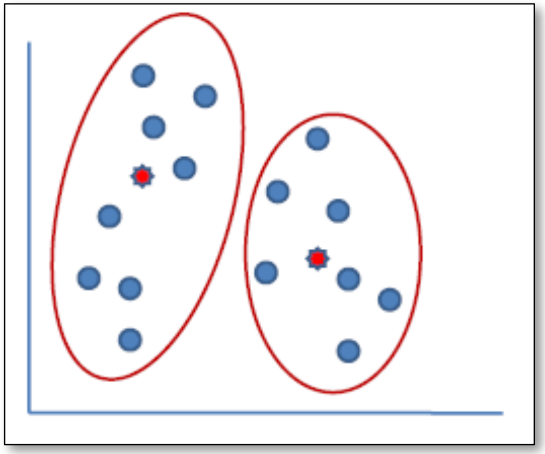

K-Means Clustering

Given set of data items and desired number of clusters *k*, K-means groups the items into *k* clusters minimizing the SSE

- Extremely difficult to compute efficiently
  - In fact, impossible
- Most algorithms compute an approximate solution (might not be absolute lowest SSE)

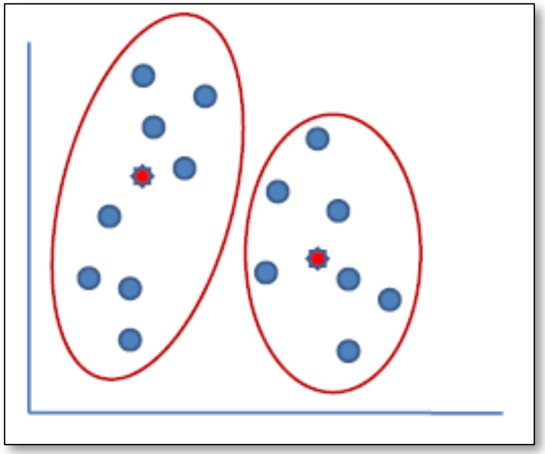

Clustering European Cities

By geographic distance, then by temperature

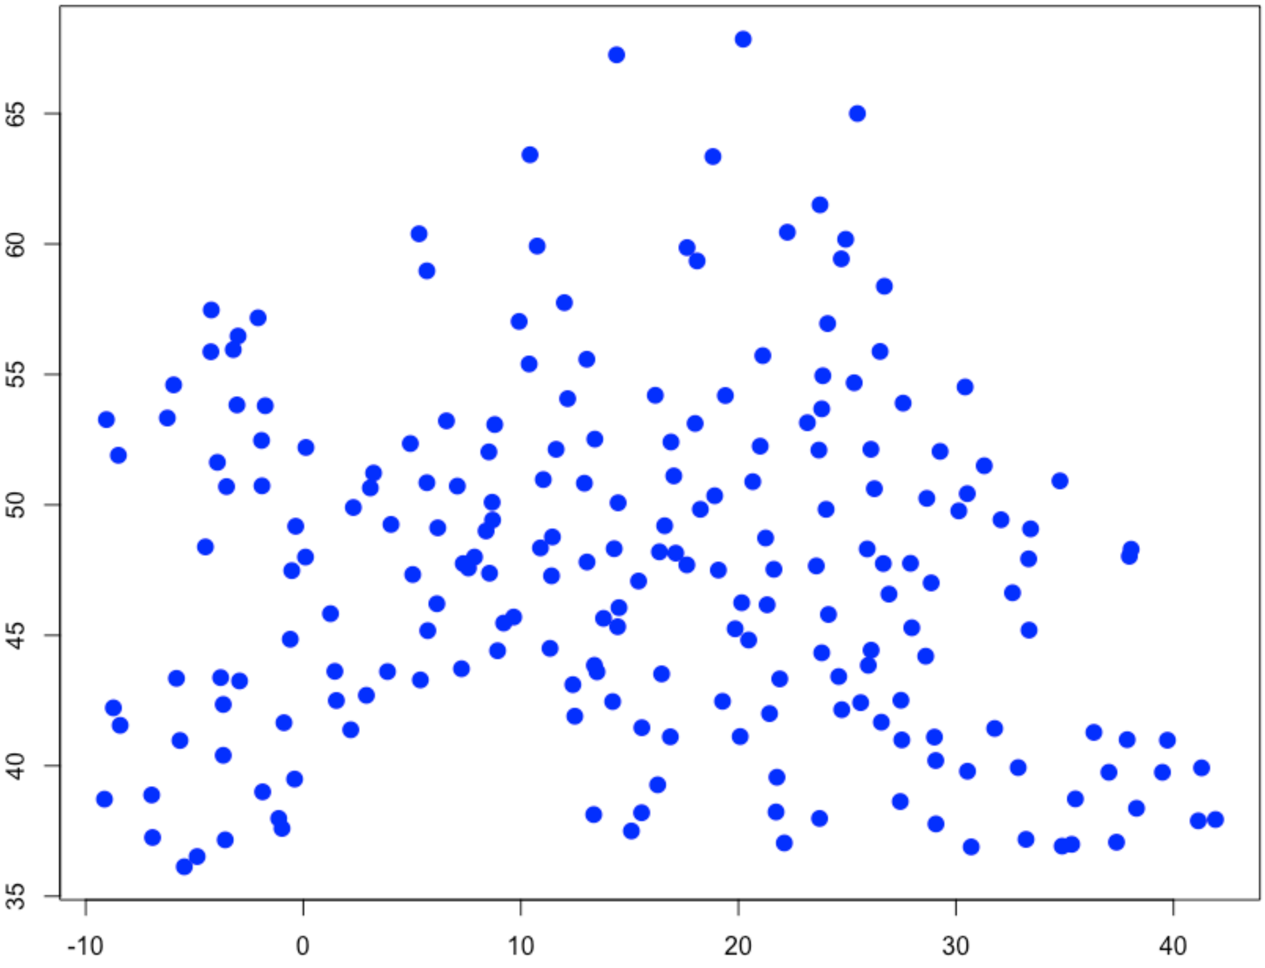

Clustering European Cities

Distance = actual distance, k = 5

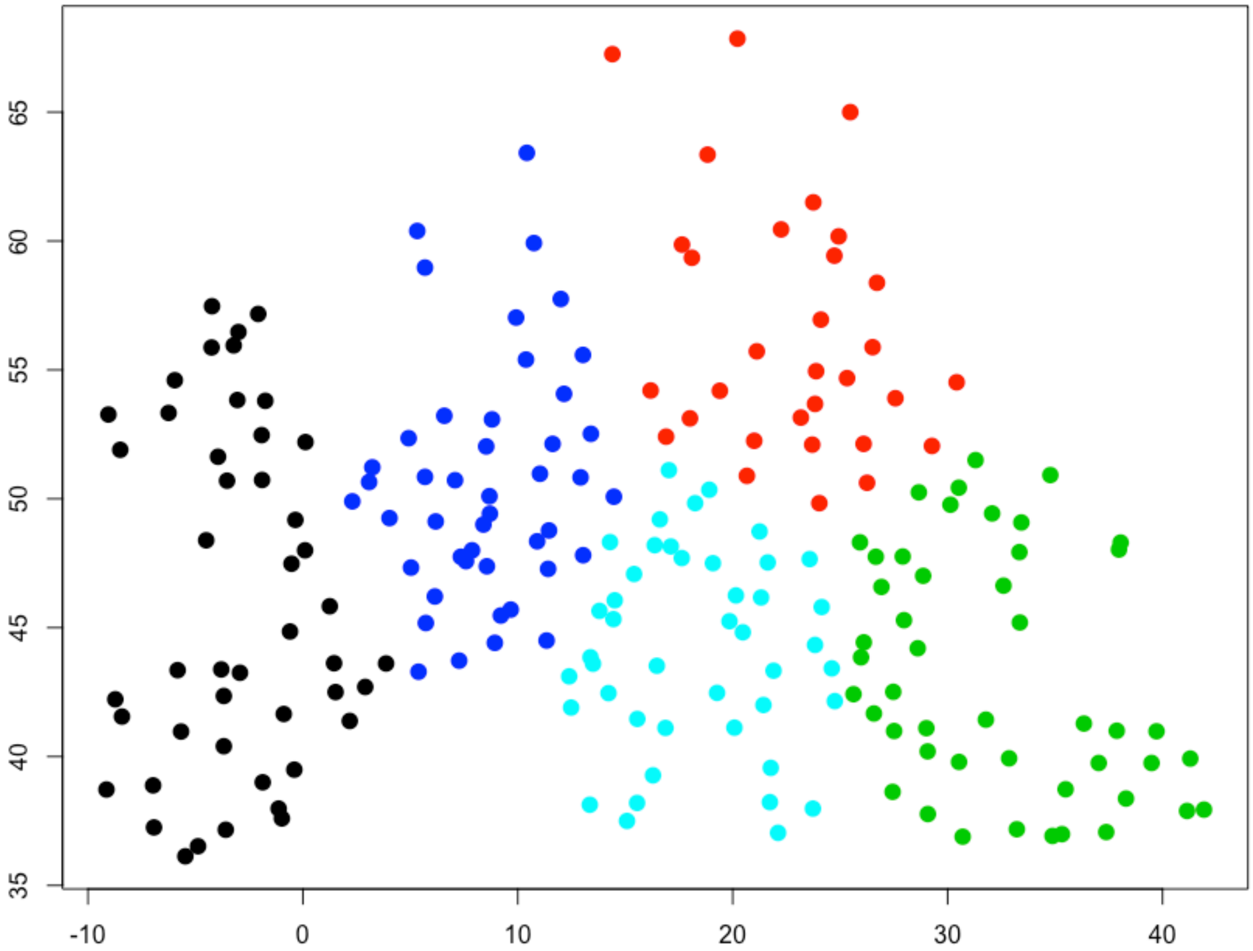

Clustering European Cities

Distance = actual distance, k = 8, with cluster means

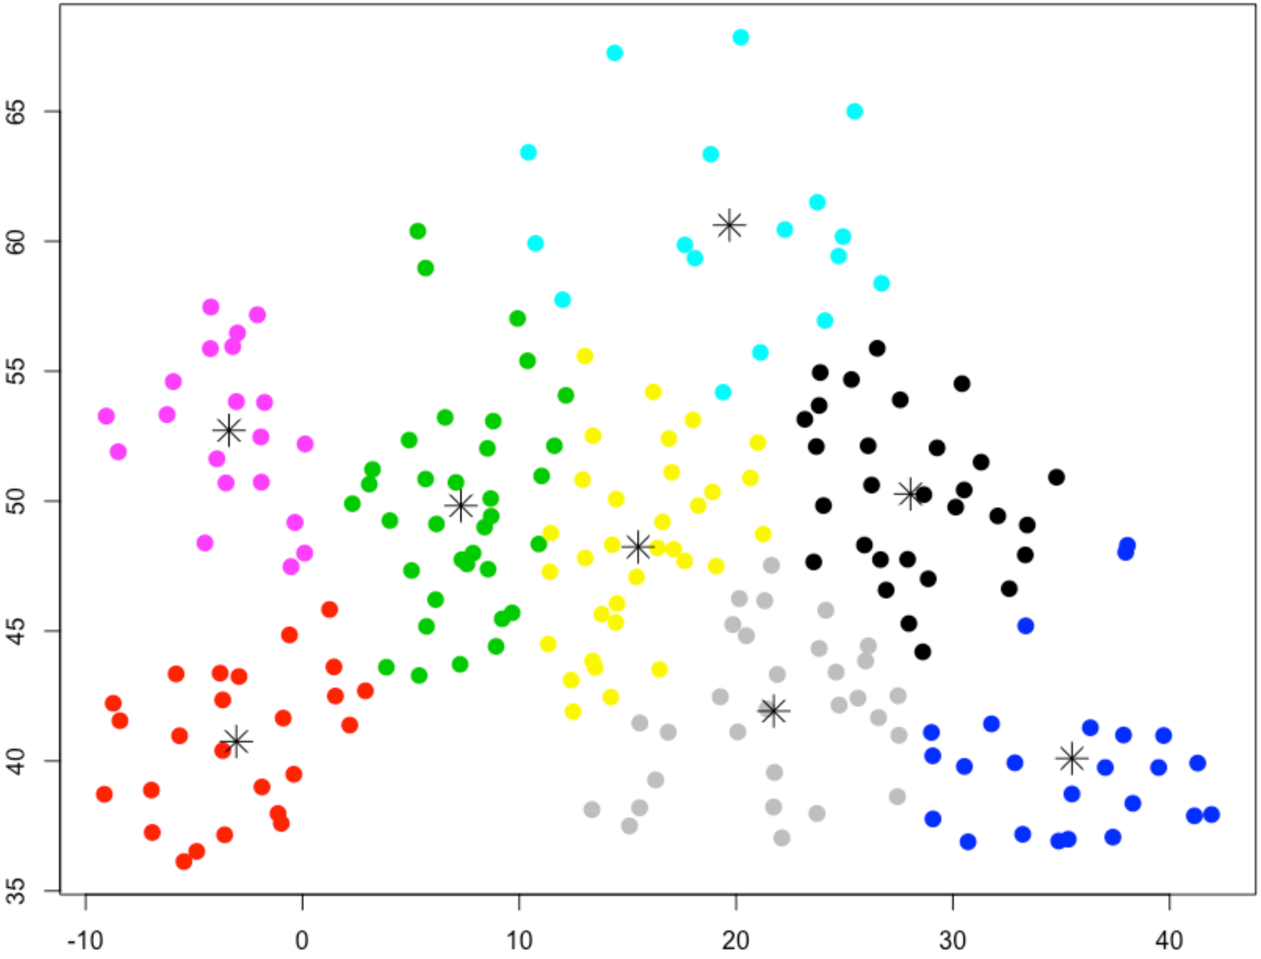

Clustering European Cities

Distance = actual distance, k = 2, with cluster means

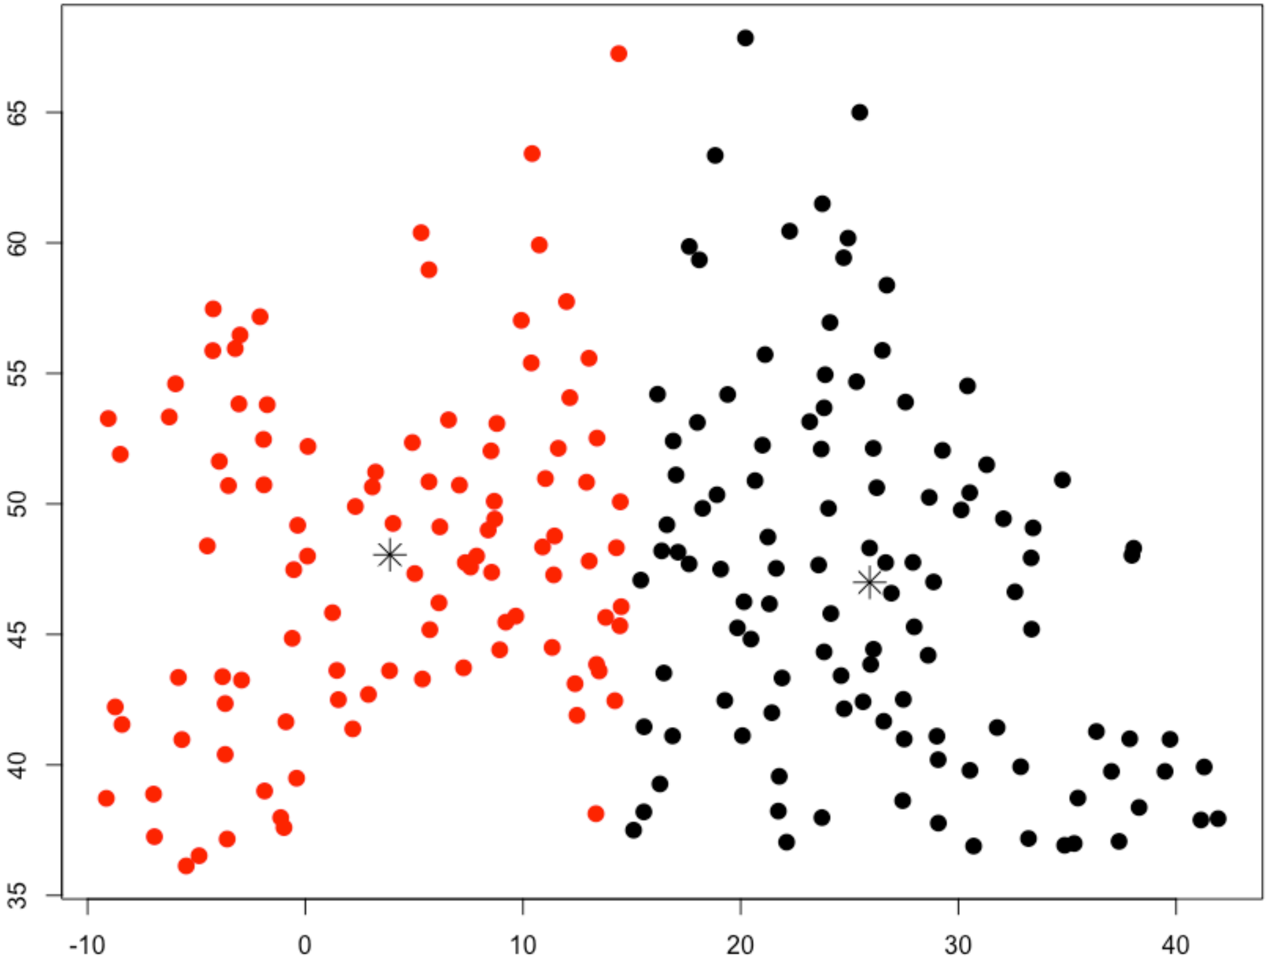

Clustering European Cities

Distance = actual distance, k = 30

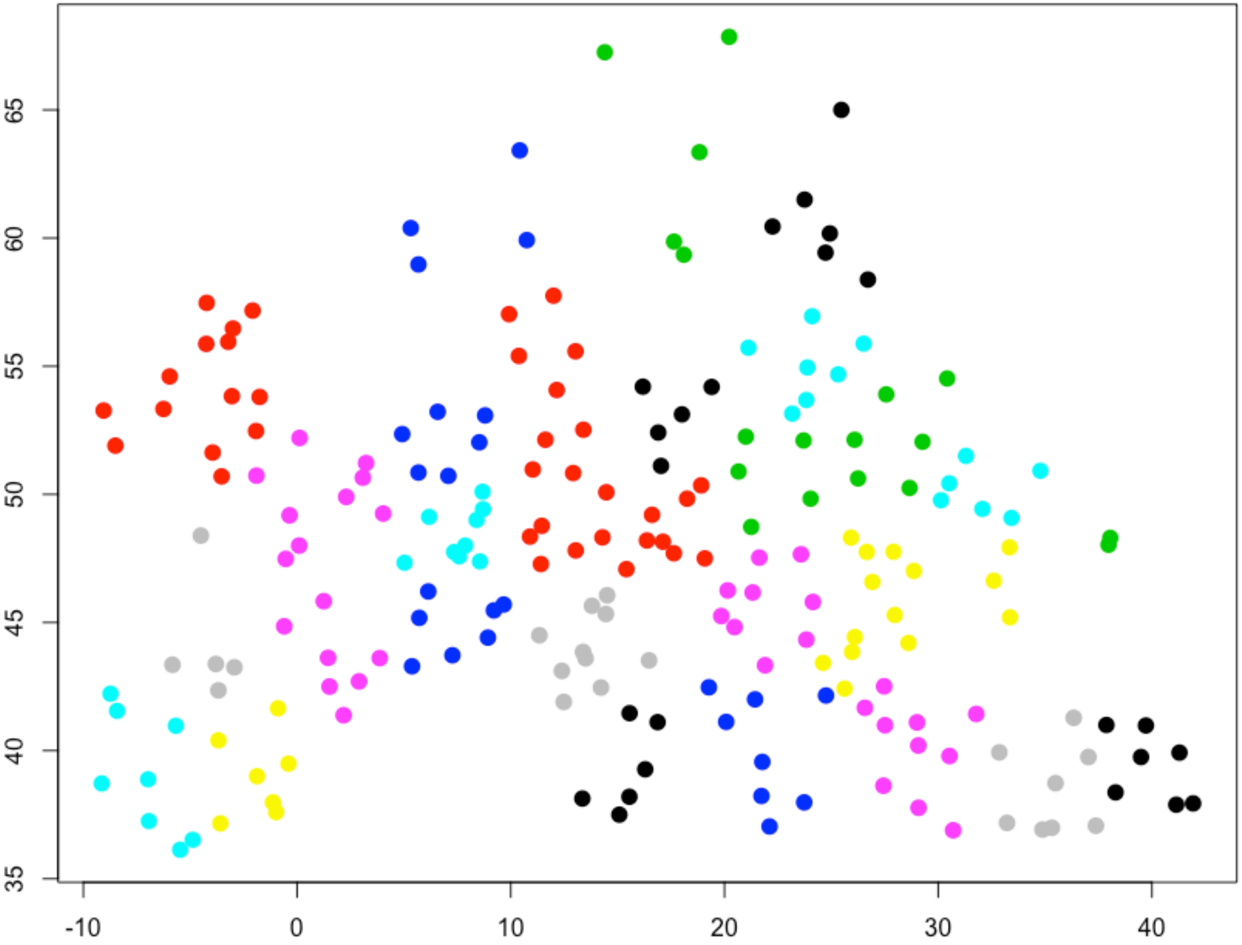

Clustering European Cities

Distance = temperature, k = 5

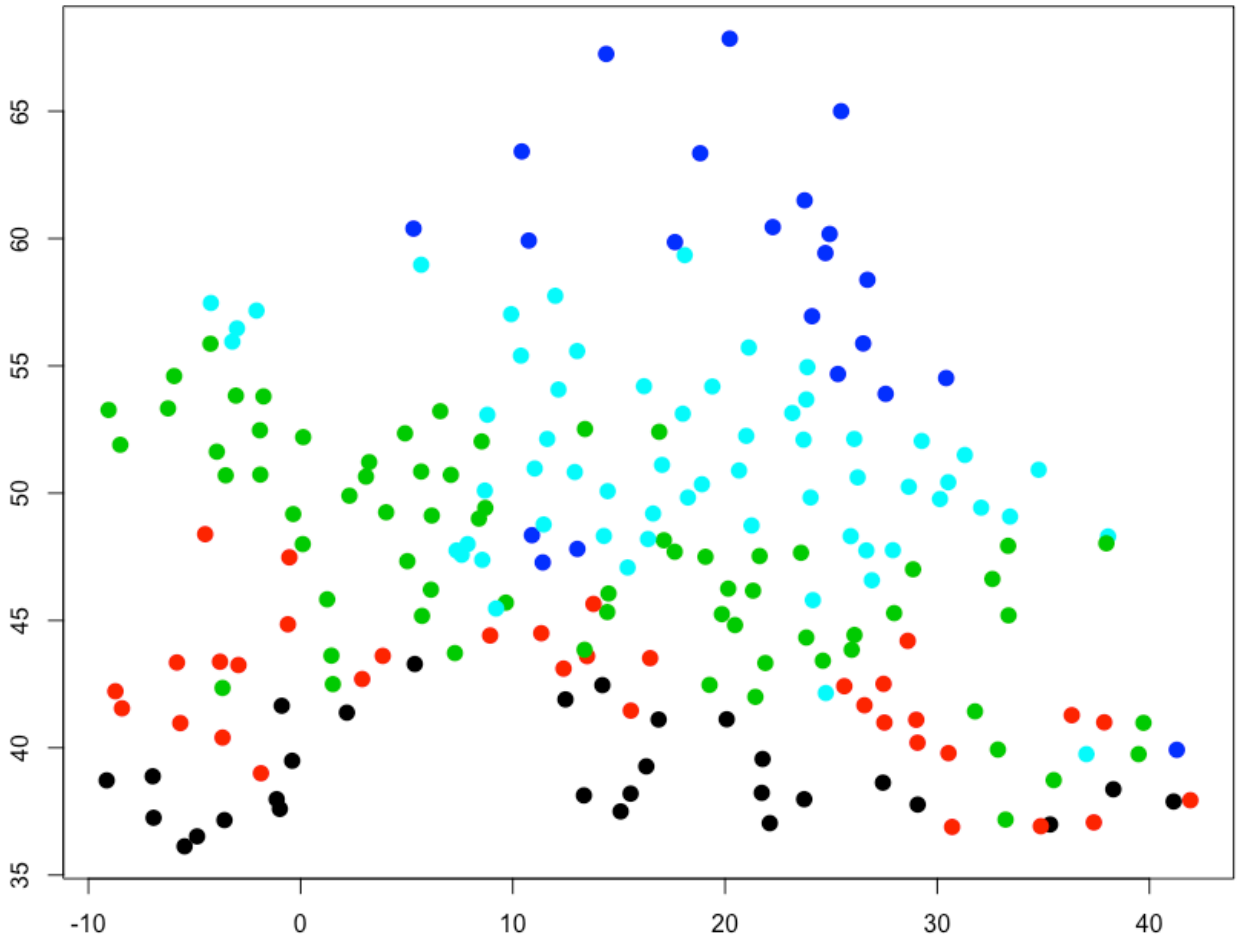

Clustering European Cities

Distance = temperature, k = 8, with means

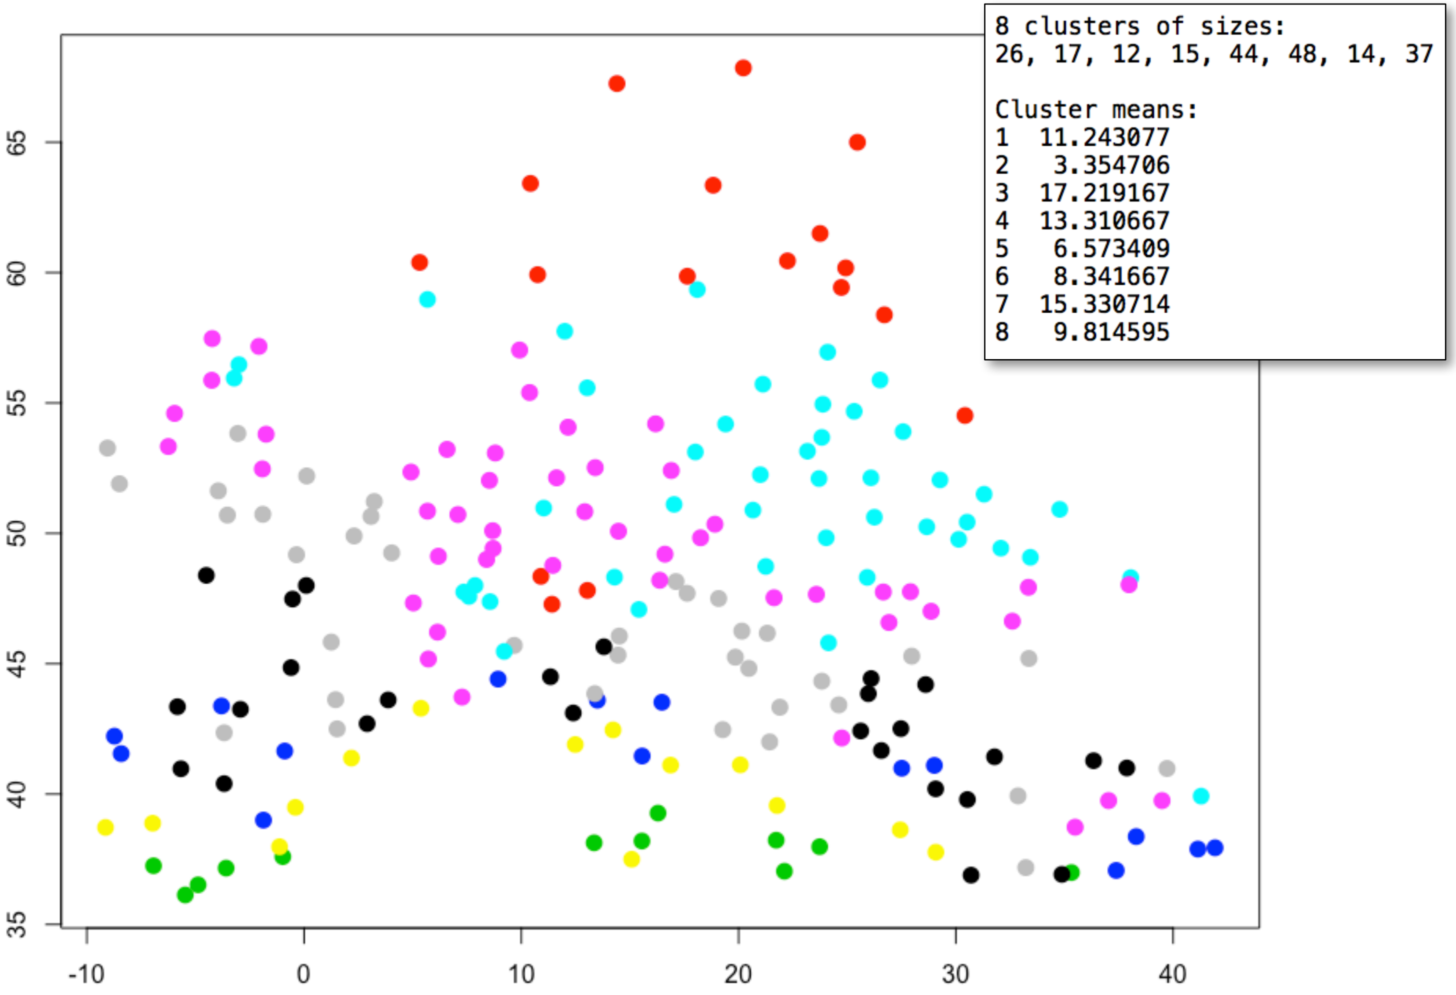

```
8 clusters of sizes:
26, 17, 12, 15, 44, 48, 14, 37

Cluster means:
1  11.243077
2   3.354706
3  17.219167
4  13.310667
5   6.573409
6   8.341667
7  15.330714
8   9.814595
```

Clustering European Cities

Distance = temperature, k = 2

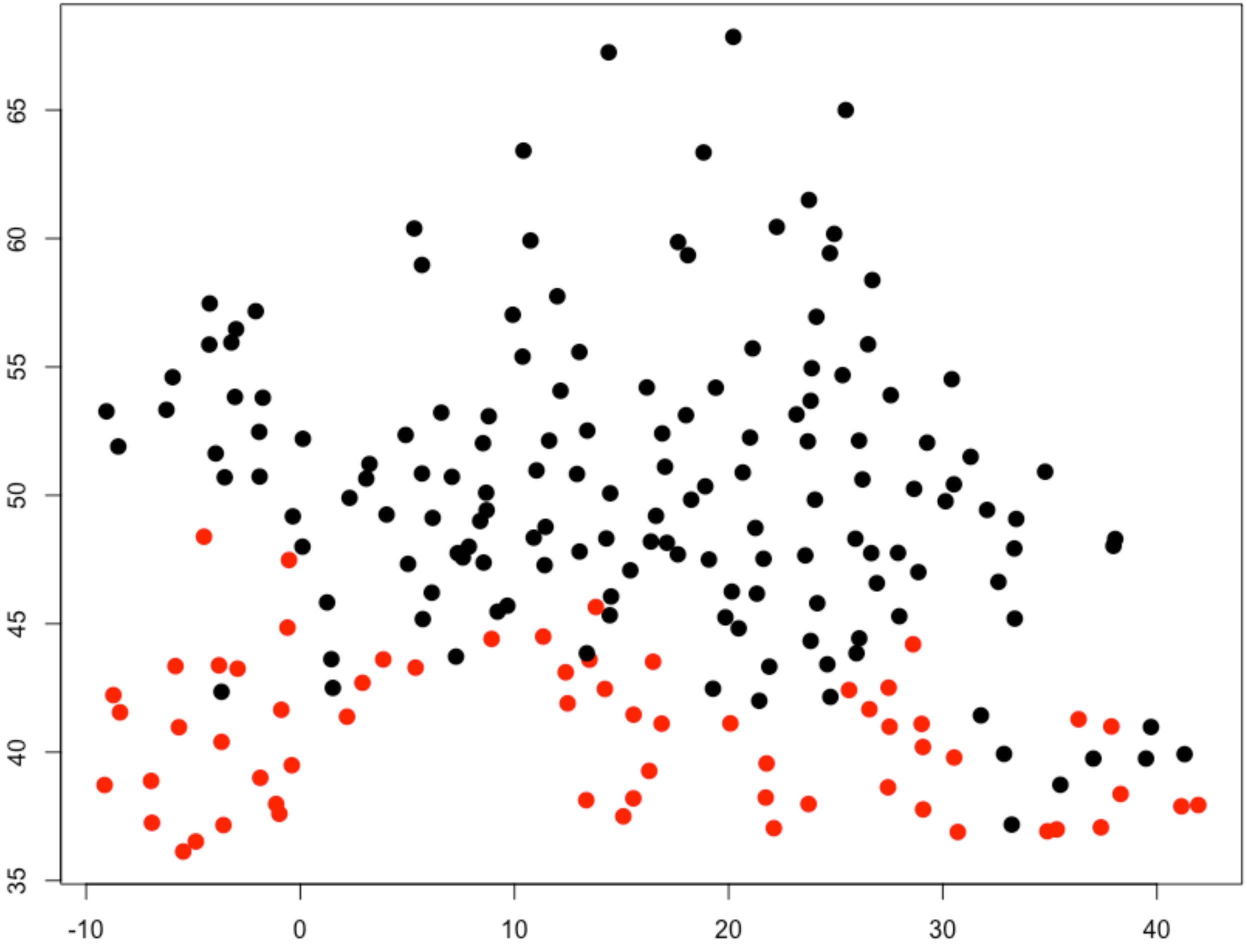

Clustering European Cities

Distance = temperature, k = 3

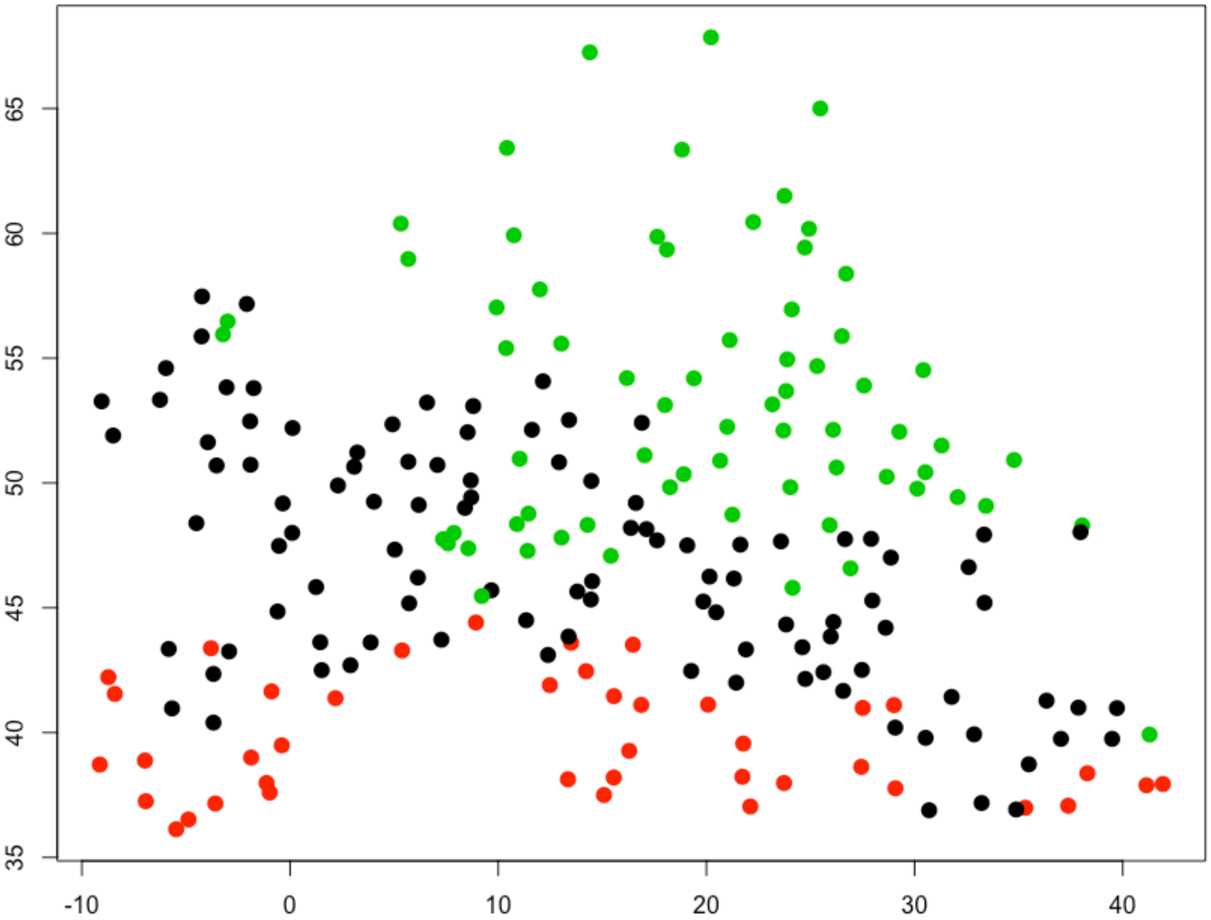

Clustering European Cities

Distance = temperature, k = 30

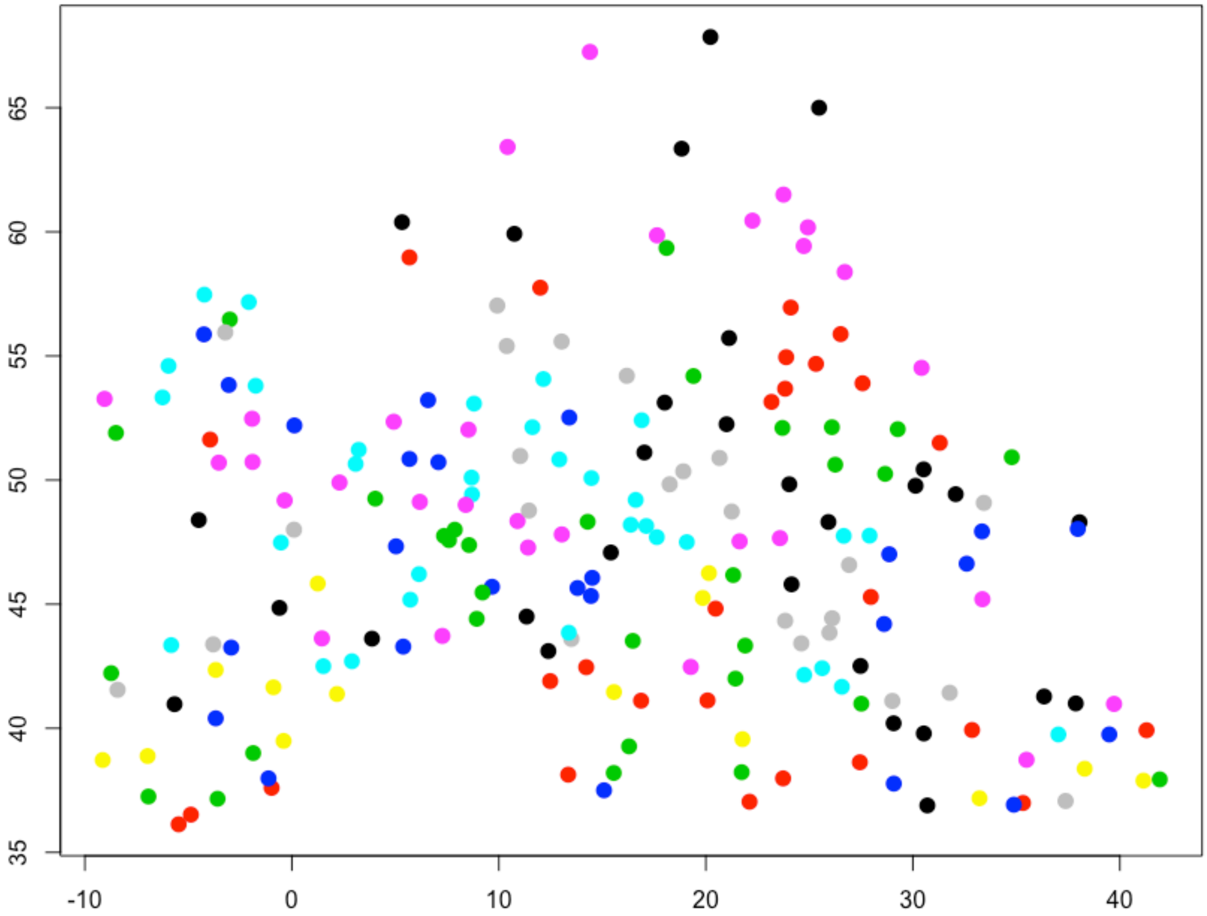

Some Uses for Clustering

- Classification
  - Assign labels to clusters
  - New data items get the label of their cluster
- Identify similar items
  - For substitutes or recommendations
  - For de-duplication
- Anomaly (outlier) detection
  - Items that are far from any cluster